

Indian Institute of Technology Jodhpur  
Operating Systems Lab (CS330)  
**Assignment 2**

Dated 17<sup>th</sup> March, 2021

Total marks: 20

This assignment is on Linux System Calls, File I/O, fork, signals, pipe etc.

This assignment has two parts. Part-I is to be completed in the lab session itself and not for evaluation. Part - II is the takeaway component and the evaluation policy is stated at the end.

**Part I:**

**(A) Write a C program that prints out all of its command-line arguments** except those that begin with a dash. Hint: find out about argc and argv. For example, after compiling your program into the file 'a.out', entering the command

a.out arg1 -arg2 arg3

prints

arg1 arg3

including a newline so your next shell prompt is not right after arg3. If all arguments begin with a dash, (or if there are no arguments at all) print nothing, that is, do not even print a newline. This is so that a blank line does not appear between the a.out command line and your next shell prompt.

**(B) Write another C program which will read** characters from its standard input stdin, count the number of NON-alphabetic ones including newlines (see the file /usr/include/ctype.h), that is, count those characters not in the a-z range nor in the A-Z range, and write out all characters read. To read and write characters, this program uses only stdin and stdout, and only the stdio library routines (see stdio.h) for input and output (see getchar and putchar). When it hits EOF in its input, the program will print out the final non-alphabetic count on stderr using fprintf and then exit(0).

**(C) You are to write yet another C program that reads** characters from stdin, reverses lowercase and uppercase letters, that is, converts to lowercase all uppercase letters and vice versa (again, see ctype.h), and writes the results onto its standard output, stdout. All characters read should be written whether converted or not. This program uses only stdin and stdout, and the stdio library routines. The program will exit(0) when it hits EOF in its input.

\*\*\*\*\*

## Hints and Notes on Linux Programming

### *Subroutines and Functions*

Declare a function to be of type void if it does not return a value, that is, if it is a subroutine rather than a true function. If you absolutely must ignore the output of a system call or function, then cast the call to void.

### *Command-line arguments*

A C main program is actually called with two arguments: argc and argv. argc is the number of command-line arguments the program is invoked with including the program name, so argc is always at least one. argv is a pointer to an array of character strings that contain the arguments, one per string.

Here is a program that echoes its arguments to show how to use these:

```
main (int argc, char *argv[])    /* echo arguments */
/* argv is not a string but an array of pointers */
/* argv[0] is command-line program name */
{
    int i;
    for (i=1; i<argc; i++)
        printf("%s%c", argv[i], (i<argc-1) ? ' ' : '\n');
}
```

Or more cryptically:

```
main (int argc, char *argv[])    /* echo arguments */
{
    while (--argc>0)
        printf((argc>1) ? "%s " : "%s\n", **++argv);
    /* argv[i] and *(argv+i) are same */
}
```

Also argv[i][j] or (\*(argv+i))[j] is the jth character of the ith command-line argument.

### *Standard input and output*

Remember IO is not a part of the C language, but is provided in the ``standard IO library'' accessed by putting

```
#include <stdio.h>
```

at the beginning of your program or in your include files. This will provide you with the functions getchar, putchar, printf, scanf and symbolic constants like EOF and NULL.

### *File access*

It is also possible to access files other than stdin and stdout directly. First, make the declaration

```
FILE *fopen(), *fp;
```

declaring the open routine and the file pointer. Then, call the open routine

```
fp = fopen(name, mode);
```

where name is a character string and mode is "r" for read, "w" for write, or "a" for append.

Now the following routines can be used for IO:

```
c = getc(fp);  
put(c, fp);
```

and also fprintf, fscanf which have a file argument.

The last thing to do is close the file (done automatically anyway at program termination) (fflush just to flush but keep the file open):

```
fclose(fp);
```

The three file pointers stdin, stdout, stderr are predefined and can be used anywhere fp above was. stderr is conventionally used for error messages like wrong arguments to a command:

```
if ((fp = fopen(++argv, "r")) == NULL) {  
    fprintf(stderr, "cat: can't open %s\n", *argv);  
    exit(1); /* close files and abort */  
} else ...
```

### *Storage allocation*

To get some dynamically allocated storage, use

```
calloc(n, sizeof(object))
```

which returns space for n objects of the given size and should be cast appropriately:

```
char *calloc();  
int *ip;  
...  
ip = (int *) calloc(n, sizeof(int));
```

Use free(p) to return the space.

## ***Linux system call interface***

These routines allow access to the Linux system at a lower, perhaps more efficient, level than the standard library, which is usually implemented in terms of what follows.

## *File descriptors*

Small positive integers are used to identify open files. Three file descriptors are automatically set up when a program is executed:

```
0 - stdin
1 - stdout
2 - stderr
```

which are usually connected with the terminal unless redirected by the shell or changed in the program with `close` and `dup`.

## *Low level I/O*

### *Opening files:*

```
int fd;
fd = open(name, rwmode);
```

where `name` is the character string file name and `rwmode` is 0, 1, 2 for read, write, read and write respectively. Better to use ```#include <fcntl.h>''` and defined constants like `O_RDONLY`. A negative result is returned in `fd` if there is an error such as trying to open a nonexistent file.

### *Creating files:*

```
fd = creat(name, pmode);
```

which creates a new file or truncates an existing file and where the octal constant `pmode` specifies the `chmod`-like 9-bit protection mode for a new file. Again a negative returned value represents an error.

### *Closing files:*

```
close(fd);
```

### *Removing files:*

```
unlink(filename);
```

### *Reading and writing files:*

```
n_read = read(fd, buf, n);
n_written = write(fd, buf, n);
```

will try to read or write `n` bytes from or to an opened or created file and return the number of bytes read or written in `n_read`, for example

```
#define BUFSIZ 1024 /* already defined by stdio.h if included earlier */
main () /* copy input to output */
{
    char buf[BUFSIZ];
```

```

    int n;
    while ((n = read(0, buf, BUFSIZ)) > 0) write(1, buf, n);
}

```

If read returns a count of 0 bytes read, then EOF has been reached.

### *Changing File Descriptors*

If you have a file descriptor fd that you want to make stdin refer to, then the following two steps will do it:

```

close(0);    /* close stdin and free slot 0 in open file table */
dup(fd);     /* dup makes a copy of fd in the smallest unused
              slot in the open file table, which is now 0 */

```

### *Standard library*

Do a ``man stdio'' to find out the standard IO library.

Do a ``man printf'', ``man getchar'', and ``man putchar'' for even more information.

Do a ``man atoi'' to find out ascii to integer conversion.

Do a ``man errno'' to find out printing error messages and the external variable errno. Also do a ``man strerror''.

Do a ``man string'' to find out the string manipulation routines available.

```

strcat      /* concatenation */
strncat
strcmp      /* compare */
strncmp
strcpy      /* copy */
strncpy
strlen      /* length */
index       /* find c in s */
rindex

```

Do a ``man 2 intro'' to find out about system calls like open, read, dup, pipe, etc.

You can do a ``man 2 open'', ``man 2 dup'', etc. to find out more about each one (``man 3 execl'', ``man 2 execve'', ``man 2 write'', ``man 2 fork'', ``man 2 close'', ``man 3 exit'', ``man 2 wait'', ``man 2 getpid'', ``man 2 alarm'', ``man 2 kill'', ``man 2 signal'').

/usr/include and /usr/include/sys contain .h files that can be included in C programs with the ``#include <stdio.h>' or ``#include <sys/inode.h>' statements, for example.

Other useful ones are ctype.h, errno.h, math.h, signal.h, string.h.

### *Error messages*

This program shows how to use perror to print error messages when system calls return an error status.

```
#include <stdio.h>
#include <sys/types.h>    /* fcntl.h needs this */
#include <fcntl.h>        /* for second argument of open */
#include <errno.h>

void exit(int);           /* gets rid of warning message */

main (int argc, char *argv[]) {
    if (open(argv[1], O_RDONLY) < 0) {
        fprintf(stderr, "errno=%d\n", errno);
        perror("open error in main");
    }
    exit(0);
}
```

### *grep and find commands*

To search all .c and .h files in the current directory for a string, do this

```
grep "search string" *.c *.h
```

To search for all files with a certain string in their name in the current directory and all its subdirectories, do this

```
find . -name "*string*" -print
```

To search all files in the current directory and all subdirectories for a certain string, do this

```
find . -exec grep "string" {} /dev/null \;
```

The /dev/null is a LINUX guru trick to get grep to print out the file name in addition to the line containing the matching string.

## Part-II

**(D) You are to combine the C programs from (A), (B) and (C) that:** (1) reads from the first file argument, reverses lowercase and uppercase letters, that is, makes lowercase all uppercase letters and vice versa, and writes out all characters, whether or not reversed; and (2) reads the above output, counts the number of NON-alphabetic characters (that is, those not in the a-z range nor in the A-Z range), and writes out all characters, whether counted or not, into its second file argument.

So that you become familiar with the various LINUX system calls and libraries (fork, execl, pipe, create, open, close, dup, stdio, exit), you will write this program in a certain contrived fashion. The program will be invoked as `driver file1 file2`. Both file arguments are required (print an error message on stderr using `fprintf` if either argument is missing and then `exit(1)`). The driver program will open the first file and create the second file, and dup them down to stdin and stdout. The driver program sets up a pipe and then forks two children.

The first child forked will dup the read end of the pipe down to stdin and then execl the program you have written, `count`, which will read characters from stdin, count the non-alphabetic ones (see `/usr/include/ctype.h`), and write all characters out to stdout. The `count` program uses only stdin and stdout, and only the stdio library routines (see `stdio.h`) for input and output (see `getchar` and `putchar`). The program will print out the final count on stderr using `fprintf` and then `exit(0)` when it hits EOF in its input.

The second child forked will dup the write end of the pipe down to stdout and then execl the program you have written, `convert`, which will read characters from stdin, reverse uppercase and lowercase (again, see `ctype.h`), and write them all out to stdout. Like `count`, this program uses only stdin and stdout. Also, it uses only the stdio library routines. The program will `exit(0)` when it hits EOF in its input.

The reason for creating the children in this order (first the one that reads from the pipe, usually called the second in the pipeline, and then second the one that writes to the pipe, usually called the first in the pipeline -- confusing!!) is that LINUX will not let a process write to a pipe if no process has the pipe open for reading; the process trying to write would get the SIGPIPE signal. So we create first the process that reads from the pipe to avoid that possibility.

Meanwhile, the parent process, the driver, will close its pipe file descriptors and then call `wait` (see `/usr/include/sys/wait.h`) twice to reap the zombie children processes when they finish. Then the parent will `exit(0)`.

Remember to close all unneeded file descriptors, including unneeded pipe ones, or EOF on a pipe read may not be detected correctly. Check all system calls for the error return (-1). For all error and debug messages, use `fprintf(stderr, ...)` or use the function `perror` or the external variable `errno` (see also `/usr/include/errno.h`). Declare a function to be of type `void` if it does not return a value, that is, if it is a subroutine rather than a true function. If you absolutely must ignore the output of a system call or function, then cast the call to `void`.

\*\*\*\*\*

## Hints and Notes

Program skeleton

- o parent opens file1:

```
#include <sys/types.h>
#include <fcntl.h>    /* defines O_RDONLY etc. */

...
fd_in = open(argv[1], O_RDONLY);    /* not fopen */
/* Check fd_in for -1!!! If it is, then the file does not
   exist or you do not have read permission. */
```

- o parent creat's file2:

```
fd_out = creat(argv[2], 0644);    /* mode = permissions, here rw-r--r--
*/
/* Check fd_out for -1!!! If it is, then the file could not
   be created */
```

- o parent uses close and dup so stdin is now file1 (done like children below)

- o parent uses close and dup so stdout is now file2: if dup returns -1, then a problem has occurred

- o parent calls pipe system call to create pipe file descriptors

- o parent forks a child to read from pipe:

- this first child manipulates file descriptors
 use close, dup so stdin is read end of pipe
 (see text Figure 1-13, similar to else clause)

- this first child execs count
 execl("count", "count", (char \*) 0)
 execl overlays the child with the binary compiled program in
 the file given by first argument and the process name becomes
 the second argument

- o parent forks again a child to write to pipe:

- this second child manipulates file descriptors
 use close, dup so stdout is write end of pipe
 (similar to if clause in text Figure 1-13)

- this second child execs convert
 execl("convert", "convert", (char \*) 0)

- o parent closes both ends of pipe



- o parent waits twice (see text Figure 1-10)

- use `wait(&status);` instead of `waitpid(-1, &status, 0);`

Before compiling and running driver, remember to:

- o in `count.c`: `fprintf(stderr, "final count = %d\n", count);`

- o `cc -o convert convert.c`, and

- o `cc -o count count.c`

The two forks form nested if statements.

```
pipe(...
if (fork() != 0) {    /* parent continues here */

    if (fork() != 0) {    /* parent continues here */
        close(write end of pipe...
        wait(...)
        wait(...)
    } else {
        ...                /* second child, writes to pipe */
    }

} else {
    ...                /* first child, reads from pipe */
}
```

It is important to check all system calls (`open`, `creat`, `dup`, etc.) for a return value `<0`, particularly `-1`, because such a return value means an error has occurred. If you just let your program continue to execute, it will just dump core at some later time and you will not know why.

It is IMPERATIVE that the parent and the first child forked (the one to read the pipe) close their write end file descriptors to the pipe. If they do not do this, then the first child will never get EOF reading from the pipe, even after the second child has sent everything. This is because EOF on a pipe is not indicated until ALL write file descriptors to the pipe are closed. If the parent and first child fail to close their file descriptors for the write end of the pipe, the program will hang forever (until killed). Furthermore, the parent must close its copy of the file descriptor to the write end of the pipe BEFORE the waits, not after, or deadlock will result. And the first child forked (the one to read the pipe) must close its file descriptor to the write end of the pipe before it starts reading the pipe, not after, or deadlock will result.

How `dup` works

`dup(fd)` looks for the smallest unused slot in the process descriptor table and makes that slot point to the same file, pipe, whatever, that `fd` points to. For example, each process starts as

process descriptor table	
-----	
slot #	points to
-----	-----
0 (stdin)	keyboard
1 (stdout)	screen
2 (stderr)	screen

If the process does a

```
fd_in = open("file1", ...
fd_out = creat("file2", ...
```

then the table now looks like

process descriptor table	
-----	
slot #	points to
-----	-----
0	keyboard
1	screen
2	screen
3	file1
4	file2

and fd\_in is 3 and fd\_out is 4. Now suppose the program does a

```
close(0);
dup(fd_in);
close(fd_in);
```

After the close(0), the table will look like

process descriptor table	
-----	
slot #	points to
-----	-----
0	-- unused --
1	screen
2	screen
3	file1
4	file2

and after the dup(fd\_in), the table will look like

process descriptor table	
-----	
slot #	points to
-----	-----
0	file1
1	screen
2	screen
3	file1

and after the `close(fd_in)`, the table will look like

```

process descriptor table
-----
slot #           points to
-----
    0             file1
    1             screen
    2             screen
    3      -- unused --
    4             file2

```

---

**(E)** Write a C program to implement a simple mindgame between a parent process and its child processes. The parent spawns two child processes C1 and C2. The parent needs to communicate with C1 and C2 on separate channel to play the game. On the other hand C1 and C2 do not need to communicate with each other. You can implement the inter-process communication through implementing any of the pipe, message queue or shared memory concept. Now, a loop runs as follows. In each iteration (also called round), P first randomly chooses one of the colors from VIBGYOR set (the rainbow color set, and the choice randomly varies from one iteration to another). Each round, both C1 and C2 also guesses one color from the VIBGYOR set independently and send to P as a string via the implemented communication channel. P reads the two strings; let these be x and y. If either of these x and y matches with the chosen color of P, then the corresponding child wins. If both the children have guessed the same color that P had chosen, then the corresponding round is ignored.

Otherwise for each round the relative distance between the color indices, denoted by *Delta*, as chosen by the parent and guessed by a child is calculated. For an example, if the parent chooses 'Green' and one of the child chooses 'Orange', then *Delta* is computed as 2. Similarly, between 'Green' and 'Indigo', the difference is 2. (We are not going to look into the direction: right or left, rather measure the relative difference in position only). For each child the score accumulates over the rounds. If the computed distance measured is x, x number of coins will be added to the corresponding child for that round. It may happen that in a round both the child earned equal number of coins, yet they did not guess the same color. If none of the child can correctly guess the same color that the parent had chosen for the round (Note the distance is 0), the child whose score crosses 50 earlier than the other, wins the game.

When one of the two children wins, the game ends, and P sends a user-defined signal to both C1 and C2. The child processes exit after handling the signal (in order to know who was the winner). After C1 and C2 exit, the parent process P exits. During each iteration of the game, P should print appropriate messages (like P's choice of the color, the strings received from C1 and

C2, which child gets the point, the current scores of C1 and C2) in order to let the user know how the game is going on.

#### Submission Instructions:

- Submit your codes in a folder as Roll.zip (Includes two c files implementing C and D component of part II respectively)
- Make sure you include a text file (Readme.txt) stating the execution status, special, conditions, assumptions or list of bugs, if any. [Your submission will not be accepted without the readme file]
- **Deadline: 23:59 hrs, 23rd March, 2021** (Any delay will deduct marks per day basis)
- For plagiarism detected in the code, '0' mark will be awarded

#### Evaluation components:

1. **Correct execution of D (10marks)**
  - Program invocation format (2)
  - Pipe implementation and management (4)
  - Implementing wait() (2)
  - Closing the file descriptor (2)
2. **Correct execution of E (10marks)**
  - Correct implementation of gaming logic (4)
  - Inter-process communication implementation (4)
  - Handling the exit strategy in the game (2)